

Unlocking 15% More Performance: A Case Study in LLVM Optimization for RISC-V

Phd Mikhail R. Gadelha
Igalia



Unlocking 15% More Performance: A Case Study in LLVM Optimization for RISC-V

Phd Mikhail R. Gadelha
Igalia



16%

Unlocking ~~15%~~ More Performance: A Case Study in LLVM Optimization for RISC-V

Phd Mikhail R. Gadelha
Igalia



Up to 16% faster SPEC CPU® 2017 on SpacemiT-X60



- This RISE project was a focused, ten-month effort to optimize the LLVM compiler for RISC-V.

Up to 16% faster SPEC CPU® 2017 on SpacemiT-X60



- This RISE project was a focused, ten-month effort to optimize the LLVM compiler for RISC-V.
- Our target board was the Banana Pi BPI-F3 with a SpacemiT-X60 8-core RISC-V processor:
 - In-order processor.
 - Supports the RVA22U64 Profile and 256-bit RVV 1.0 standard.

Up to 16% faster SPEC CPU® 2017 on SpacemiT-X60



- This RISE project was a focused, ten-month effort to optimize the LLVM compiler for RISC-V.
- Our target board was the Banana Pi BPI-F3 with a SpacemiT-X60 8-core RISC-V processor:
 - In-order processor.
 - Supports the RVA22U64 Profile and 256-bit RVV 1.0 standard.
- Our goal: to close the performance gap between LLVM and the GCC compiler.

Up to 16% faster SPEC CPU® 2017 on SpacemiT-X60



- This RISE project was a focused, ten-month effort to optimize the LLVM compiler for RISC-V.
- Our target board was the Banana Pi BPI-F3 with a SpacemiT-X60 8-core RISC-V processor:
 - In-order processor.
 - Supports the RVA22U64 Profile and 256-bit RVV 1.0 standard.
- Our goal: to close the performance gap between LLVM and the GCC compiler.
- Our result: individual contributions boosted performance by up to 16% on SPEC CPU® 2017 benchmarks.

The Project

- Prior to this work, a clear performance gap existed between code generated by LLVM and GCC for RISC-V.

The Project

- Prior to this work, a clear performance gap existed between code generated by LLVM and GCC for RISC-V.
- There is no single solution to close the gap, as improvements and regressions occur daily within the codebase.

The Project

- Prior to this work, a clear performance gap existed between code generated by LLVM and GCC for RISC-V.
- There is no single solution to close the gap, as improvements and regressions occur daily within the codebase.
- This presentation will focus on our three main contributions to help close the gap:
 - Introducing a scheduling model for the SpacemiT-X60.
 - Improvements to vectorization across calls.
 - Register Allocation with IPRA Support for RISC-V.

Our Contributions

(major contributions first)

SpacemiT-X60 Scheduling Model

SpacemiT-X60 Scheduling Model

- Instruction Scheduling == Performance.
- Wrong latencies/resources → compiler makes poor choices.



```
1 fld      ft0, 0(a0)
2 fadd.d   ft1, ft0, ft2
3 fmul.d   ft3, ft4, ft5
```

before




```
1 fld      ft0, 0(a0)
2 fmul.d   ft3, ft4, ft5
3 fadd.d   ft1, ft0, ft2
```

after


SpacemiT-X60 Scheduling Model

- Instruction Scheduling == Performance.
- Wrong latencies/resources → compiler makes poor choices.



SpacemiT-X60 Scheduling Model

- Instruction Scheduling == Performance.
- Wrong latencies/resources → compiler makes poor choices.



SpacemiT-X60 Scheduling Model

- Instruction Scheduling == Performance.
- Wrong latencies/resources → compiler makes poor choices.

```
1 fld ft0, 0(a0)
2 fadd.d ft1, ft0, ft2
3 fmul.d ft3, ft4, ft5
```

Issued later,
independent


before

```
1 fld ft0, 0(a0)
2 fmul.d ft3, ft4, ft5
3 fadd.d ft1, ft0, ft2
```

after


SpacemiT-X60 Scheduling Model

- Instruction Scheduling == Performance.
- Wrong latencies/resources → compiler makes poor choices.




SpacemiT-X60 Scheduling Model

- Instruction Scheduling == Performance.
- Wrong latencies/resources → compiler makes poor choices.



SpacemiT-X60 Scheduling Model

- Instruction Scheduling == Performance.
- Wrong latencies/resources → compiler makes poor choices.



How We Got the Numbers

- We built custom microbenchmarks to measure instruction latencies.


How We Got the Numbers

- We built custom microbenchmarks to measure instruction latencies.
- Most of instruction throughput data available at https://camel-cdr.github.io/rvv-bench-results/bpi_f3/index.html.


How We Got the Numbers

- We built custom microbenchmarks to measure instruction latencies.
- Most of instruction throughput data available at https://camel-cdr.github.io/rvv-bench-results/bpi_f3/index.html.
- It's RISC but:
 - 201 scalar instructions.
 - 82 floating-point instructions.
 - 9185 RVV instructions (because of the combination of different LMULs and SEWs).


RVA22U64 SPEC exec time, O3+LTO+mcpu=spacemit-x60




RVA22U64 SPEC exec time, O3+LTO+mcpu=spacemit-x60




RVA22U64_V SPEC exec time, O3+LTO+mcpu=spacemit-x60



RVA22U64_V SPEC exec time, O3+LTO+mcpu=spacemit-x60




RVA22U64 vs RVA22U64_V, O3+LTO+mcpu=spacemit-x60



- Scheduling nearly eliminated the gap between scalar and vector configs.

RVA22U64 vs RVA22U64_V, O3+LTO+mcpu=spacemit-x60



- Scheduling nearly eliminated the gap between scalar and vector configs.
- On in-order processors like X60, scheduling is critical; on out-of-order, impact would be smaller and vectorization more decisive.

Improvements to Vectorization Across Calls

Improvements to Vectorization Across Calls



- Initial surprise: vectorized code sometimes underperformed scalar.

Improvements to Vectorization Across Calls


- Initial surprise: vectorized code sometimes underperformed scalar.
- Root cause: poor cost modeling and suboptimal spill behavior.

Improvements to Vectorization Across Calls




- Initial surprise: vectorized code sometimes underperformed scalar.
- Root cause: poor cost modeling and suboptimal spill behavior.
- The extra cycles were due to register spilling, particularly around function call boundaries.

The 544.nab_r case




```
define void @f(i1 %c, ptr %p, ptr %q) {
entry:
    %x0 = load i64, ptr %p
    %p1 = getelementptr i64, ptr %p, i64 1
    %x1 = load i64, ptr %p1
    br i1 %c, label %foo, label %bar
foo:
    call void @g()
    br label %baz
bar:
    call void @g()
    br label %baz
baz:
    store i64 %x0, ptr %q
    %q1 = getelementptr i64, ptr %q, i64 1
    store i64 %x1, ptr %q1
    ret void
}
```



The 544.nab_r case

Loads first
value from %p

```
define void @f(i1 %c, ptr %p, ptr %q) {
entry:
    %x0 = load i64, ptr %p
    %p1 = getelementptr i64, ptr %p, i64 1
    %x1 = load i64, ptr %p1
    br i1 %c, label %foo, label %bar
foo:
    call void @g()
    br label %baz
bar:
    call void @g()
    br label %baz
baz:
    store i64 %x0, ptr %q
    %q1 = getelementptr i64, ptr %q, i64 1
    store i64 %x1, ptr %q1
    ret void
}
```




The 544.nab_r case

Loads second value from %p

```


1 define void @f(i1 %c, ptr %p, ptr %q) {
2 entry:
3     %x0 = load i64, ptr %p
4     %p1 = getelementptr i64, ptr %p, i64 1
5     %x1 = load i64, ptr %p1
6     br i1 %c, label %foo, label %bar
7 foo:
8     call void @g()
9     br label %baz
10 bar:
11    call void @g()
12    br label %baz
13 baz:
14    store i64 %x0, ptr %q
15    %q1 = getelementptr i64, ptr %q, i64 1
16    store i64 %x1, ptr %q1
17    ret void
18 }
```



The 544.nab_r case

Conditional
jump


```
1 define void @f(i1 %c, ptr %p, ptr %q) {
2 entry:
3     %x0 = load i64, ptr %p
4     %p1 = getelementptr i64, ptr %p, i64 1
5     %x1 = load i64, ptr %p1
6     br i1 %c, label %foo, label %bar
7 foo:
8     call void @g()
9     br label %baz
10 bar:
11    call void @g()
12    br label %baz
13 baz:
14    store i64 %x0, ptr %q
15    %q1 = getelementptr i64, ptr %q, i64 1
16    store i64 %x1, ptr %q1
17    ret void
18 }
```



The 544.nab_r case

Conditional call
to g()


```
1 define void @f(i1 %c, ptr %p, ptr %q) {
2 entry:
3     %x0 = load i64, ptr %p
4     %p1 = getelementptr i64, ptr %p, i64 1
5     %x1 = load i64, ptr %p1
6     br i1 %c, label %foo, label %bar
7 foo:
8     call void @g()
9     br label %baz
10 bar:
11    call void @g()
12    br label %baz
13 baz:
14    store i64 %x0, ptr %q
15    %q1 = getelementptr i64, ptr %q, i64 1
16    store i64 %x1, ptr %q1
17    ret void
18 }
```




The 544.nab_r case

Conditional call
to g()


```
1 define void @f(i1 %c, ptr %p, ptr %q) {
2 entry:
3     %x0 = load i64, ptr %p
4     %p1 = getelementptr i64, ptr %p, i64 1
5     %x1 = load i64, ptr %p1
6     br i1 %c, label %foo, label %bar
7 foo:
8     call void @g()
9     br label %baz
10 bar:
11     call void @g()
12     br label %baz
13 baz:
14     store i64 %x0, ptr %q
15     %q1 = getelementptr i64, ptr %q, i64 1
16     store i64 %x1, ptr %q1
17     ret void
18 }
```



The 544.nab_r case




```
1 define void @f(i1 %c, ptr %p, ptr %q) {
2 entry:
3     %x0 = load i64, ptr %p
4     %p1 = getelementptr i64, ptr %p, i64 1
5     %x1 = load i64, ptr %p1
6     br i1 %c, label %foo, label %bar
7 foo:
8     call void @g()
9     br label %baz
10 bar:
11    call void @g()
12    br label %baz
13 baz:
14    store i64 %x0, ptr %q
15    %q1 = getelementptr i64, ptr %q, i64 1
16    store i64 %x1, ptr %q1
17    ret void
18 }
```



The 544.nab_r case

Stores second
value from %q

```
1 define void @f(i1 %c, ptr %p, ptr %q) {
2 entry:
3     %x0 = load i64, ptr %p
4     %p1 = getelementptr i64, ptr %p, i64 1
5     %x1 = load i64, ptr %p1
6     br i1 %c, label %foo, label %bar
7 foo:
8     call void @g()
9     br label %baz
10 bar:
11    call void @g()
12    br label %baz
13 baz:
14     store i64 %x0, ptr %q
15     %q1 = getelementptr i64, ptr %q, i64 1
16     store i64 %x1, ptr %q1
17     ret void
18 }
```



Fixing Real Bugs

- Since we are storing and loading from continuous regions, the accesses can be vectorized.


Fixing Real Bugs

- Since we are storing and loading from continuous regions, the accesses can be vectorized.
- We found that the SLP Vectorizer was aggressively vectorizing regions without properly accounting for the cost of spilling vector registers across calls.

Fixing Real Bugs

- Since we are storing and loading from continuous regions, the accesses can be vectorized.
- We found that the SLP Vectorizer was aggressively vectorizing regions without properly accounting for the cost of spilling vector registers across calls.
- Previously, the SLP vectorizer only analyzed the entry and baz blocks, ignoring foo and bar entirely.

The 544.nab_r case




```

1 define void @f(i1 %c, ptr %p, ptr %q) {
2 entry:
3     %x0 = load i64, ptr %p
4     %p1 = getelementptr i64, ptr %p, i64 1
5     %x1 = load i64, ptr %p1
6     br i1 %c, label %foo, label %bar
7     foo:
8         call void @g()
9         br label %baz
10    bar:
11        call void @g()
12        br label %baz
13    baz:
14        store i64 %x0, ptr %q
15        %q1 = getelementptr i64, ptr %q, i64 1
16        store i64 %x1, ptr %q1
17        ret void
18 }

```

Not being analyzed



Addressing the Issue

- To address the issue, we first proposed a patch which modified the SLP vectorizer to properly walk through all basic blocks when analyzing cost.

Addressing the Issue

- To address the issue, we first proposed a patch which modified the SLP vectorizer to properly walk through all basic blocks when analyzing cost.
- Promising results: execution time dropped by **9.92%** in 544.nab_r.


Addressing the Issue

- To address the issue, we first proposed a patch which modified the SLP vectorizer to properly walk through all basic blocks when analyzing cost.
- Promising results: execution time dropped by **9.92%** in 544.nab_r.
- But with a major drawback: **+6.9%** increase compilation time in 502.gcc_r.


Addressing the Issue

- To address the issue, we first proposed a patch which modified the SLP vectorizer to properly walk through all basic blocks when analyzing cost.
- Promising results: execution time dropped by **9.92%** in `544.nab_r`.
- But with a major drawback: **+6.9%** increase compilation time in `502.gcc_r`.
- Following discussions with the community, Alexey Bataev (SLP Vectorizer code owner) proposed and landed refined solution.

RVA22U64_V SPEC exec time, O3+LTO, SLP fix



RVA22U64_V SPEC exec time, O3+LTO, SLP fix



IPRA (Inter-Procedural Register Allocation) Support

IPRA (Inter-Procedural Register Allocation) Support



- Function calls can often spill (when you store a register to the stack) more registers than necessary → wasted cycles on every call.

IPRA (Inter-Procedural Register Allocation) Support



- Function calls can often spill (when you store a register to the stack) more registers than necessary → wasted cycles on every call.
- IPRA: caller/callee register use is tracked across the functions, by eliminating unnecessary save/restore sequences.

What IPRA Brings



```
1 foo:  
2     addi    sp, sp, -32  
3     sd      ra, 24(sp)  
4     sd      s0, 16(sp)  
5     sd      s1, 8(sp)  
6     ...  
7     ld      s1, 8(sp)  
8     ld      s0, 16(sp)  
9     ld      ra, 24(sp)  
10    addi   sp, sp, 32  
11    ret
```

before



```
1 foo:  
2     addi    sp, sp, -8  
3     sd      ra, 0(sp)  
4     ...  
5     ld      ra, 0(sp)  
6     addi   sp, sp, 8  
7     ret
```

after

What IPRA Brings

before

```

1 foo:
2   addi    sp, sp, -3
3   sd      ra, 24(sp)
4   sd      s0, 16(sp)
5   sd      s1, 8(sp)
6   ...
7   ld      s1, 8(sp)
8   ld      s0, 16(sp)
9   ld      ra, 24(sp)
10  addi   sp, sp, 32
11  ret

```

after

```

1 foo:
2   addi    sp, sp, -8
3   sd      ra, 0(sp)
4   ...
5   ld      ra, 0(sp)
6   addi   sp, sp, 8
7   ret

```

saved even if
not really live

saved even if
not really live

What IPRA Brings

```
1 foo:  
2     addi    sp, sp, -32  
3     sd      ra, 24(sp)  
4     sd      s0, 16(sp)  
5     sd      s1, 8(sp)  
6     ...  
7     ld      s1, 8(sp)  
8     ld      s0, 16(sp)  
9     ld      ra, 24(sp)  
10    addi   sp, sp, 32  
11    ret
```

before

```
1 foo:  
2     addi    sp, sp, -8  
3     sd      ra, 0(sp)  
4     ...  
5     ld      ra, 0(sp)  
6     addi   sp, sp, 8  
7     ret
```

after

only what's
truly needed

Impact of IPRA

- Reduction in register pressure, shorter prologue/epilogue code.


Impact of IPRA

- Reduction in register pressure, shorter prologue/epilogue code.
- SPEC benchmarks showed measurable improvements (small but consistent).


Impact of IPRA

- Reduction in register pressure, shorter prologue/epilogue code.
- SPEC benchmarks showed measurable improvements (small but consistent).
- Unfortunately, it can't be enabled by default: IPRA is not enabled by default due to a bug (described in issue [119556](#)), however, it does not affect the SPEC benchmarks.


RVA22U64 SPEC exec time, O3+LTO+IPRA




RVA22U64 SPEC exec time, O3+LTO+IPRA



RVA22U64_V SPEC exec time, O3+LTO+IPRA



RVA22U64_V SPEC exec time, O3+LTO+IPRA



Conclusions

Faster RISC-V Today

- We contributed with:
 - Scheduling: largest wins, especially for scalar-heavy code.
 - Vectorization: enabled smarter spilling cost calculations.
 - IPRA: smaller but consistent improvements across workloads.

Faster RISC-V Today

- We contributed with:
 - Scheduling: largest wins, especially for scalar-heavy code.
 - Vectorization: enabled smarter spilling cost calculations.
 - IPRA: smaller but consistent improvements across workloads.
- Almost all changes are upstream benefiting everyone. Under review:
 - <https://github.com/llvm/llvm-project/pull/150618>
 - <https://github.com/llvm/llvm-project/pull/150644>
 - <https://github.com/llvm/llvm-project/pull/152557>
 - <https://github.com/llvm/llvm-project/pull/152737>
 - <https://github.com/llvm/llvm-project/pull/152738>

What We Learned Along the Way

- Scheduling is critical for performance.
 - No scheduling model → LLVM pessimises the final code.
 - We should likely adopt some scheduling model as default, like other targets do.
 - Should we make the X60 scheduling model default for in-order RISC-V processors?

What We Learned Along the Way

- Scheduling is critical for performance.
 - No scheduling model → LLVM pessimises the final code.
 - We should likely adopt some scheduling model as default, like other targets do.
 - Should we make the X60 scheduling model default for in-order RISC-V processors?
- Many contributions don't have immediate benchmark impact.

What We Learned Along the Way

- Scheduling is critical for performance.
 - No scheduling model → LLVM pessimizes the final code.
 - We should likely adopt some scheduling model as default, like other targets do.
 - Should we make the X60 scheduling model default for in-order RISC-V processors?
- Many contributions don't have immediate benchmark impact.
- Vectorization needs careful tuning to avoid regressions.

Did we close the performance gap between LLVM
and the GCC compiler?

LLVM vs GCC on SpacemIT-X60

- Note it's not a direct apples-to-apples comparison.

LLVM vs GCC on SpacemiT-X60

- Note it's not a direct apples-to-apples comparison.
- The code was compiled with the same RISC-V extensions enabled.


LLVM vs GCC on SpacemiT-X60

- Note it's not a direct apples-to-apples comparison.
- The code was compiled with the same RISC-V extensions enabled.
- But GCC doesn't have X60-specific scheduling latencies, while LLVM does.


LLVM vs GCC on SpacemiT-X60

- Note it's not a direct apples-to-apples comparison.
- The code was compiled with the same RISC-V extensions enabled.
- But GCC doesn't have X60-specific scheduling latencies, while LLVM does.
- Still useful to show relative progress and identify where LLVM has caught up.

RVA22U64 SPEC execution time, GCC vs LLVM



RVA22U64_V SPEC execution time, GCC vs LLVM



Thanks!

- This work at Igalia was made possible thanks to support from RISE, under Project RP009.
- Thanks to my Igalia colleagues for discussions, and feedback.
- And to the LLVM RISC-V community for reviews and for getting patches upstream quickly.

- Accidental Dataflow Analysis: Extending the RISC-V VL Optimizer @ 2025 EuroLLVM by Luke Lau:
<https://www.youtube.com/watch?v=bkOwPr36SrQ>
- Improvements to RISC-V Vector code generation in LLVM @ 2025 RISC-V Summit Europe by Luke Lau and Alex Bradbury:
<https://www.youtube.com/watch?v=0NjugW7FF48>
- RISC-V nightly performance testing of top-of-tree GCC and clang:
<https://cc-perf.igalia.com/>



NORTH AMERICA 2025

